

## Financial Risk Analysis with Python(Goldman Sachs Financial Analysis Dataset)

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv('goldman_sachs.csv')
df.head()
```

```
Out[2]:
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product	Firm	Region	Manager	TransactionDate	TransactionAmount	AccountBalance	RiskScore	CreditRating	TenureMo
0	33	CUST6549	ACC12334	Credit	Withdrawal	Savings Account	Firm C	Central	Manager 1	21-10-2023	87480.05448	74008.43310	0.729101	319	
1	177	CUST2942	ACC52650	Credit	Withdrawal	Home Loan	Firm A	East	Manager 3	20-06-2023	20315.74505	22715.83590	0.472424	692	
2	178	CUST6776	ACC45101	Current	Deposit	Personal Loan	Firm C	South	Manager 3	02-01-2023	10484.57165	42706.09210	0.648784	543	
3	173	CUST2539	ACC88252	Current	Withdrawal	Mutual Fund	Firm A	Central	Manager 2	25-07-2023	45122.27373	114176.56870	0.734832	430	
4	67	CUST2626	ACC21878	Savings	Withdrawal	Home Loan	Firm C	Central	Manager 4	25-07-2023	42360.79878	17863.02644	0.289304	468	

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   TransactionID    800 non-null    int64  
 1   CustomerID       800 non-null    object  
 2   AccountID        800 non-null    object  
 3   AccountType      800 non-null    object  
 4   TransactionType   800 non-null    object  
 5   Product          800 non-null    object  
 6   Firm              800 non-null    object  
 7   Region            800 non-null    object  
 8   Manager           800 non-null    object  
 9   TransactionDate  800 non-null    object  
 10  TransactionAmount 800 non-null    float64 
 11  AccountBalance   800 non-null    float64 
 12  RiskScore         800 non-null    float64 
 13  CreditRating     800 non-null    int64  
 14  TenureMonths     800 non-null    int64  
dtypes: float64(3), int64(3), object(9)
memory usage: 93.9+ KB
```

```
In [4]: df.isnull().sum()
```

```
Out[4]:
```

TransactionID	0
CustomerID	0
AccountID	0
AccountType	0
TransactionType	0
Product	0
Firm	0
Region	0
Manager	0
TransactionDate	0
TransactionAmount	0
AccountBalance	0
RiskScore	0
CreditRating	0
TenureMonths	0
dtype:	int64

```
In [5]: for col in df.columns:
    if df[col].nunique()<10:
        print(df[col].value_counts())
        print("----*40)
```

```
AccountType
Loan      218
Credit    206
Savings   204
Current   172
Name: count, dtype: int64
```

```
-----
```

```
TransactionType
Withdrawal 207
Payment    200
Deposit    199
Transfer   194
Name: count, dtype: int64
```

```
-----
```

```
Product
Home Loan    176
Credit Card   168
Mutual Fund   162
Personal Loan 153
Savings Account 141
Name: count, dtype: int64
```

```
-----
```

```
Firm
Firm B      168
Firm C      165
Firm D      164
Firm E      153
Firm A      150
Name: count, dtype: int64
```

```
-----
```

```
Region
West       176
East       162
South      158
North      156
Central    148
Name: count, dtype: int64
```

```
-----
```

```
Manager
Manager 1   218
Manager 3   209
Manager 2   193
Manager 4   180
Name: count, dtype: int64
```

```
In [6]: df.duplicated().sum()

# If duplicates exist and you want to drop:
df = df.drop_duplicates()
df.shape
```

```
Out[6]: (800, 15)
```

#### Task 1: Data Cleaning and Formatting

- Remove/treat any special characters or non-numeric entries from financial fields.
- Convert currency amounts into numerical format.
- Validate and format date columns.
- Ensure account types and transaction categories are standardized.

```
In [7]: # type 1.1 ● Remove/treat any special characters or non-numeric entries from financial fields.
df["TransactionAmount"] = df["TransactionAmount"].abs()

financial_cols = ["TransactionAmount", "AccountBalance"]
for col in financial_cols:
    df[col] = pd.to_numeric(df[col], errors="coerce")
df.columns
df.head()
```

```
Out[7]: TransactionID CustomerID AccountID AccountType TransactionType Product Firm Region Manager TransactionDate TransactionAmount AccountBalance RiskScore CreditRating TenureMo
0 33 CUST6549 ACC12334 Credit Withdrawal Savings Account Firm C Central Manager 1 21-10-2023 87480.05448 74008.43310 0.729101 319
1 177 CUST2942 ACC52650 Credit Withdrawal Home Loan Firm A East Manager 3 20-06-2023 20315.74505 22715.83590 0.472424 692
2 178 CUST6776 ACC45101 Current Deposit Personal Loan Firm C South Manager 3 02-01-2023 10484.57165 42706.09210 0.648784 543
3 173 CUST2539 ACC88252 Current Withdrawal Mutual Fund Firm A Central Manager 2 25-07-2023 45122.27373 114176.56870 0.734832 430
4 67 CUST2626 ACC21878 Savings Withdrawal Home Loan Firm C Central Manager 4 25-07-2023 42360.79878 17863.02644 0.289304 468
```

```
In [8]: # step 1.2 ● Convert currency amounts into numerical format.
df.columns = df.columns.str.strip()

text_cols = df.select_dtypes(include="object").columns
df[text_cols] = df[text_cols].apply(lambda s: s.str.strip())
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
 # Column      Non-Null Count Dtype  
--- 
0 TransactionID 800 non-null   int64  
1 CustomerID   800 non-null   object  
2 AccountID    800 non-null   object  
3 AccountType   800 non-null   object  
4 TransactionType 800 non-null   object  
5 Product       800 non-null   object  
6 Firm          800 non-null   object  
7 Region        800 non-null   object  
8 Manager        800 non-null   object  
9 TransactionDate 800 non-null   object  
10 TransactionAmount 800 non-null   float64 
11 AccountBalance 800 non-null   float64 
12 RiskScore     800 non-null   float64 
13 CreditRating  800 non-null   int64  
14 TenureMonths  800 non-null   int64  
dtypes: float64(3), int64(3), object(9)
memory usage: 93.9+ KB
```

```
In [9]: # step 1.3 ● Validate and format date columns.
df["TransactionDate"] = df["TransactionDate"].apply(pd.to_datetime, format="%d-%m-%Y")
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
 # Column      Non-Null Count Dtype  
--- 
0 TransactionID 800 non-null   int64  
1 CustomerID   800 non-null   object  
2 AccountID    800 non-null   object  
3 AccountType   800 non-null   object  
4 TransactionType 800 non-null   object  
5 Product       800 non-null   object  
6 Firm          800 non-null   object  
7 Region        800 non-null   object  
8 Manager        800 non-null   object  
9 TransactionDate 800 non-null   datetime64[ns]
10 TransactionAmount 800 non-null   float64 
11 AccountBalance 800 non-null   float64 
12 RiskScore     800 non-null   float64 
13 CreditRating  800 non-null   int64  
14 TenureMonths  800 non-null   int64  
dtypes: datetime64[ns](1), float64(3), int64(3), object(8)
memory usage: 93.9+ KB
```

```
In [10]: # step 1.4 ● Ensure account types and transaction categories are standardized.

df["AccountType"] = df["AccountType"].str.strip().str.title()
df["TransactionType"] = df["TransactionType"].str.strip().str.title()
df[["TransactionType", "AccountType"]]
```

```
Out[10]:
```

	TransactionType	AccountType
0	Withdrawal	Credit
1	Withdrawal	Credit
2	Deposit	Current
3	Withdrawal	Current
4	Withdrawal	Savings
...	...	...
795	Withdrawal	Current
796	Withdrawal	Credit
797	Payment	Current
798	Payment	Current
799	Deposit	Loan

800 rows × 2 columns

#### Task 2: Descriptive Transactional Analysis

- Calculate monthly and yearly summaries of total credits, debits, and net transaction volume.
- Plot trends in total credits vs. debits over time.
- Identify top and bottom performing accounts based on net inflow.
- Identify and flag accounts as dormant or inactive if there is a gap of two months or more between consecutive transactions.

```
In [11]: # 2.1 ● Calculate monthly and yearly summaries of total credits, debits, and net transaction volume.
credit_keywords = {"deposit"}
debit_keywords = {"withdrawal", "transfer", "payment"}

#Separate Credit and Debit amounts

df["Credit"] = np.where(
    df["TransactionType"].str.title() == "deposit",
    df["TransactionAmount"],
    0
)

df["Debit"] = np.where(
    df["TransactionType"].str.title() == "withdrawal",
    df["TransactionAmount"],
    0
)

df[["TransactionType", "TransactionAmount", "Credit", "Debit"]].head()

def txn_sign(txn_type: str) -> int:
    t = str(txn_type).lower()
    if any(k in t for k in debit_keywords):
        return -1
    if any(k in t for k in credit_keywords):
        return 1
    return 1 # default if unknown (or set 0)

df["Sign"] = df["TransactionType"].apply(txn_sign)
df["NetAmount"] = df["TransactionAmount"] * df["Sign"]

df["Year"] = df["TransactionDate"].dt.year
df["Month"] = df["TransactionDate"].dt.to_period("M").astype(str)

monthly = df.groupby("Month").agg(
    TotalCredits=("TransactionAmount", lambda s: s[df.loc[s.index, "Sign"]==1].sum()),
    TotalDebits=("TransactionAmount", lambda s: s[df.loc[s.index, "Sign"]==-1].sum()),
    Net=("NetAmount","sum")
).reset_index()

yearly = df.groupby("Year").agg(
    TotalCredits=("TransactionAmount", lambda s: s[df.loc[s.index, "Sign"]==1].sum()),
    TotalDebits=("TransactionAmount", lambda s: s[df.loc[s.index, "Sign"]==-1].sum()),
    Net=("NetAmount","sum")
).reset_index()

df.head()
```

```
Out[11]:
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product	Firm	Region	Manager	TransactionDate	...	AccountBalance	RiskScore	CreditRating	TenureMonths	Credit	Del
0	33	CUST6549	ACC12334	Credit	Withdrawal	Savings Account	Firm C	Central	Manager 1	2023-10-21	...	74008.43310	0.729101	319	200	0.0	0
1	177	CUST2942	ACC52650	Credit	Withdrawal	Home Loan	Firm A	East	Manager 3	2023-06-20	...	22715.83590	0.472424	692	47	0.0	0
2	178	CUST6776	ACC45101	Current	Deposit	Personal Loan	Firm C	South	Manager 3	2023-01-02	...	42706.09210	0.648784	543	109	0.0	0
3	173	CUST2539	ACC88252	Current	Withdrawal	Mutual Fund	Firm A	Central	Manager 2	2023-07-25	...	114176.56870	0.734832	430	103	0.0	0
4	67	CUST2626	ACC21878	Savings	Withdrawal	Home Loan	Firm C	Central	Manager 4	2023-07-25	...	17863.02644	0.289304	468	234	0.0	0

5 rows × 21 columns

```
In [12]: import seaborn as sns
import matplotlib.pyplot as plt

# Force white background everywhere
sns.set_theme(style="whitegrid")
plt.rcParams["figure.facecolor"] = "white"
plt.rcParams["axes.facecolor"] = "white"

plt.figure(figsize=(15, 7))

# Credits line
sns.lineplot(
    data=monthly,
    x="Month",
    y="TotalCredits",
    marker="o",
    linestyle="--",
    label="Credits",
    color="purple"
)
```

```

# Debits line
sns.lineplot(
    data=monthly,
    x="Month",
    y="TotalDebits",
    marker="o",
    linestyle="--",
    label="Debits",
    color="coral"
)

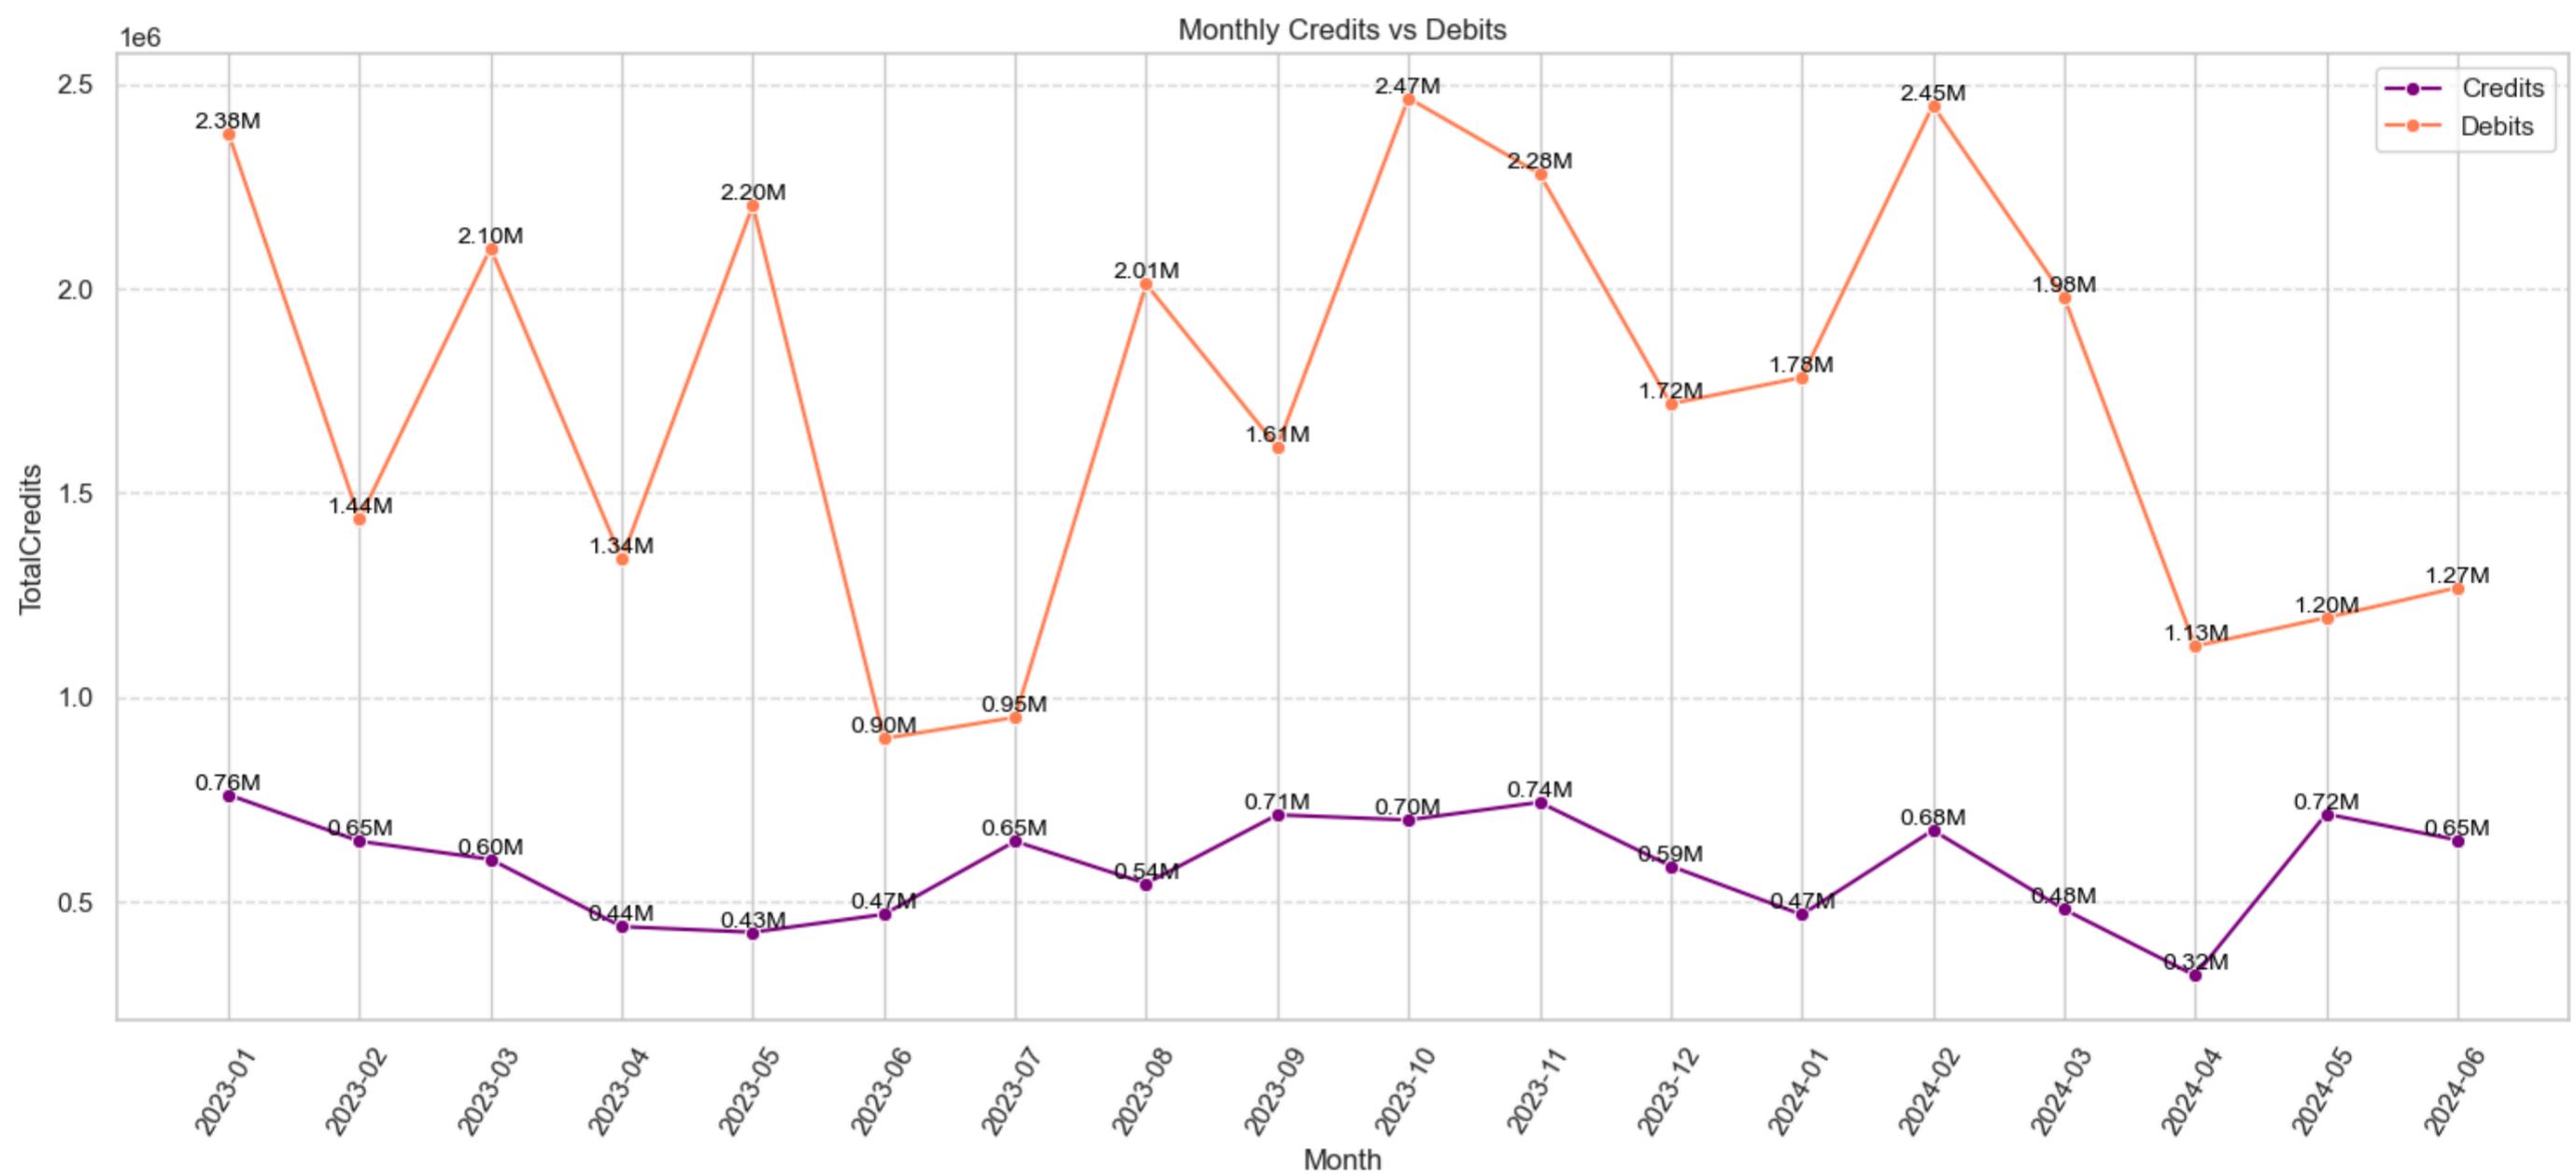
plt.xticks(rotation=60)
plt.grid(axis="y", linestyle="--", alpha=0.6)
plt.title("Monthly Credits vs Debits")
plt.legend()

# Credit labels
for x, y in zip(monthly["Month"], monthly["TotalCredits"]):
    plt.text(
        x, y,
        f"{y/1e6:.2f}M",
        ha="center",
        va="bottom",
        fontsize=10,
        color="black"
    )

# Debit labels
for x, y in zip(monthly["Month"], monthly["TotalDebits"]):
    plt.text(
        x, y,
        f"{y/1e6:.2f}M",
        ha="center",
        va="bottom",
        fontsize=10,
        color="black"
    )

plt.tight_layout()
plt.show()

```



```

In [13]: # step 2.3 • Identify top and bottom performing accounts based on net inflow.
# Identify Top Performing Accounts
top_accounts = df.sort_values(
    by='NetAmount',
    ascending=False
).head(5)

# Identify Bottom Performing Accounts
bottom_accounts = df.sort_values(
    by='NetAmount',
    ascending=True
).head(5)

print(f"\nTop Performing Accounts:\n{top_accounts[['AccountBalance', 'NetAmount']]}\n")

print(f"\nBottom Performing Accounts:\n{bottom_accounts[['AccountBalance', 'NetAmount']]}\n")

```

Top Performing Accounts:

	AccountBalance	NetAmount
741	51092.98942	130149.3799
237	33532.61157	127718.9013
234	93754.78157	122856.4830
379	42088.94404	118019.7086
269	32952.55399	113221.7353

Bottom Performing Accounts:

	AccountBalance	NetAmount
790	116386.890400	-130726.9141
80	14174.176850	-130475.1663
338	42085.195500	-129983.6051
319	150199.078600	-128816.2360
320	9228.815058	-127907.5731

```

In [14]: # step 2.4 • Identify and flag accounts as dormant or inactive if there is a gap of two months or more between consecutive transactions.
df_sorted = df.sort_values(["AccountID", "TransactionDate"])
df_sorted["PrevDate"] = df_sorted.groupby("AccountID")["TransactionDate"].shift(1)
df_sorted["GapDays"] = (df_sorted["TransactionDate"] - df_sorted["PrevDate"]).dt.days

# ~60 days as 2 months (simple rule)
df_sorted["DormantFlag"] = df_sorted["GapDays"] >= 60

dormant_accounts = (df_sorted.groupby("AccountID")["DormantFlag"].any()
                     .reset_index(name="IsDormant"))

df.head()

```

Out[14]:	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product	Firm	Region	Manager	TransactionDate	...	AccountBalance	RiskScore	CreditRating	TenureMonths	Credit	De
0	33	CUST6549	ACC12334	Credit	Withdrawal	Savings Account	Firm C	Central	Manager 1	2023-10-21	...	74008.43310	0.729101	319	200	0.0	(
1	177	CUST2942	ACC52650	Credit	Withdrawal	Home Loan	Firm A	East	Manager 3	2023-06-20	...	22715.83590	0.472424	692	47	0.0	)
2	178	CUST6776	ACC45101	Current	Deposit	Personal Loan	Firm C	South	Manager 3	2023-01-02	...	42706.09210	0.648784	543	109	0.0	(
3	173	CUST2539	ACC88252	Current	Withdrawal	Mutual Fund	Firm A	Central	Manager 2	2023-07-25	...	114176.56870	0.734832	430	103	0.0	)
4	67	CUST2626	ACC21878	Savings	Withdrawal	Home Loan	Firm C	Central	Manager 4	2023-07-25	...	17863.02644	0.289304	468	234	0.0	(

5 rows × 21 columns

### Task 3: Customer Profile Building

- Group accounts by activity levels: High, Medium, Low based on transaction frequency on your analysis and rubrics. Do not forget to mention the rubric in the headings.
- Segment customers by average balance and transaction volume.
- Create profiles for:
  - High-net inflow accounts
  - High-frequency low-balance accounts
  - Accounts with negative or near-zero balances

```
In [15]: # step 3.1 ● Group accounts by activity levels: High, Medium, Low based on transaction frequency on your
#           analysis and rubrics. Do not forget to mention the rubric in the headings.

txn_freq = (df.groupby(["CustomerID", "Month"])
            .size()
            .reset_index(name="TxnCount"))

avg_monthly_txn = txn_freq.groupby("CustomerID")["TxnCount"].mean().reset_index(name="AvgMonthlyTxn")

print(avg_monthly_txn)
```

CustomerID	AvgMonthlyTxn
0 CUST1042	1.142857
1 CUST1114	1.333333
2 CUST1121	1.333333
3 CUST1189	1.000000
4 CUST1223	1.000000
.. ...	...
183 CUST9683	1.750000
184 CUST9731	1.500000
185 CUST9754	1.333333
186 CUST9843	1.000000
187 CUST9962	1.666667

[188 rows × 2 columns]

```
In [16]: df["CustomerID"].nunique()
```

Out[16]: 188

```
In [17]: txn_freq = (df.groupby(["CustomerID", "Month"])
                 .value_counts()
                 .reset_index(name="TxnCount"))

avg_monthly_txn = txn_freq.groupby("CustomerID")["TxnCount"].mean().reset_index(name="AvgMonthlyTxn")

def activity_level(x):
    if x >= 20: return "High"
    if x >= 5: return "Medium"
    return "Low"

avg_monthly_txn["ActivityLevel"] = avg_monthly_txn["AvgMonthlyTxn"].apply(activity_level)

print(avg_monthly_txn["ActivityLevel"].tail(100))
```

88 Low
89 Low
90 Low
91 Low
92 Low
...
183 Low
184 Low
185 Low
186 Low
187 Low

Name: ActivityLevel, Length: 100, dtype: object

```
In [18]: # step 3.2 ● Segment customers by average balance and transaction volume.

account_metrics = (
    df.groupby("AccountID")
    .agg(
        Total_Credit=("Credit", "sum"),
        Total_Debit=("Debit", "sum"),
        Transaction_Volume=("TransactionAmount", "sum")
    )
    .assign(Net_Inflow=lambda x: x["Total_Credit"] - x["Total_Debit"])
    .reset_index()
)

account_metrics.head()
```

AccountID	Total_Credit	Total_Debit	Transaction_Volume	Net_Inflow
0 ACC10117	0.0	0.0	199480.967430	0.0
1 ACC10996	0.0	0.0	250739.550950	0.0
2 ACC11062	0.0	0.0	27189.136160	0.0
3 ACC11188	0.0	0.0	257576.603590	0.0
4 ACC11285	0.0	0.0	96729.609841	0.0

```
In [19]: # step 3.3 ● Create profiles for:
```

# a:) ○ High-net inflow accounts:

```
high_inflow = account_metrics["Net_Inflow"].quantile(0.75)
```

```

low_inflow = account_metrics["Net_Inflow"].quantile(0.25)
high_volume = account_metrics["Transaction_Volume"].quantile(0.75)
low_volume = account_metrics["Transaction_Volume"].quantile(0.25)

def segment_customer(row):
    if row["Net_Inflow"] >= high_inflow or row["Transaction_Volume"] >= high_volume:
        return "High Value"
    elif row["Net_Inflow"] <= low_inflow or row["Transaction_Volume"] <= low_volume:
        return "Low Value"
    else:
        return "Medium Value"

account_metrics["Customer_Segment"] = account_metrics.apply(segment_customer, axis=1)
account_metrics.head()

```

Out[19]:

	AccountID	Total_Credit	Total_Debit	Transaction_Volume	Net_Inflow	Customer_Segment
0	ACC10117	0.0	0.0	199480.967430	0.0	High Value
1	ACC10996	0.0	0.0	250739.550950	0.0	High Value
2	ACC11062	0.0	0.0	27189.136160	0.0	High Value
3	ACC11188	0.0	0.0	257576.603590	0.0	High Value
4	ACC11285	0.0	0.0	96729.609841	0.0	High Value

In [20]:

```

account_metrics = (
    df.groupby("AccountID")
    .agg(
        avg_balance=("AccountBalance", "mean"),
        transaction_volume=("TransactionAmount", "sum")
    )
    .reset_index()
)
account_metrics.head()

```

Out[20]:

	AccountID	avg_balance	transaction_volume
0	ACC10117	70107.007957	199480.967430
1	ACC10996	43568.008084	250739.550950
2	ACC11062	38137.132610	27189.136160
3	ACC11188	69652.151044	257576.603590
4	ACC11285	97401.348560	96729.609841

In [21]:

```

# b:) ○ High-frequency low-balance accounts:

cust_summary = df.groupby("CustomerID").agg(
    AvgBalance=("AccountBalance","mean"),
    TotalVolume=("TransactionAmount","sum"),
    TotalNet=("NetAmount","sum"),
    TxnCount=("TransactionID","count")
).reset_index()

# Simple quantile segmentation

cust_summary["BalanceSegment"] = pd.qcut(cust_summary["AvgBalance"], 3, labels=["LowBal","MidBal","HighBal"])
cust_summary["VolumeSegment"] = pd.qcut(cust_summary["TotalVolume"], 3, labels=["LowVol","MidVol","HighVol"])

print(cust_summary["BalanceSegment"], cust_summary["VolumeSegment"])

```

```

0      HighBal
1      LowBal
2      HighBal
3      LowBal
4      HighBal
...
183     MidBal
184     HighBal
185     HighBal
186     MidBal
187     HighBal
Name: BalanceSegment, Length: 188, dtype: category
Categories (3, object): ['LowBal' < 'MidBal' < 'HighBal'] 0      HighVol
1      HighVol
2      HighVol
3      LowVol
4      LowVol
...
183     HighVol
184     HighVol
185     MidVol
186     MidVol
187     HighVol
Name: VolumeSegment, Length: 188, dtype: category
Categories (3, object): ['LowVol' < 'MidVol' < 'HighVol']

```

In [22]:

```

#c:) ○ Accounts with negative or near-zero balances:

high_net_inflow = cust_summary.sort_values("TotalNet", ascending=False).head(20)

high_freq_low_bal = cust_summary[
    (cust_summary["TxnCount"] >= cust_summary["TxnCount"].quantile(0.75)) &
    (cust_summary["AvgBalance"] <= cust_summary["AvgBalance"].quantile(0.25))
]

neg_or_near_zero = cust_summary[cust_summary["AvgBalance"] <= 0]

print(high_freq_low_bal)

```

```

CustomerID AvgBalance TotalVolume TotalNet TxnCount \
1 CUST1114 60127.893845 385802.565919 -268029.670739 8
53 CUST059 59259.390588 296413.616370 -296413.616370 6
54 CUST3069 48347.423620 179440.309010 -97487.457390 5
80 CUST4584 52584.244758 298235.343843 -64988.794803 5
94 CUST5253 47802.857312 177843.520130 -157496.948350 5
99 CUST5428 58148.843446 261052.014510 -212246.224430 8
101 CUST5545 61151.359039 265333.196444 -265333.196444 7
109 CUST5912 61966.849233 354409.131240 -172474.809920 6
111 CUST5920 61623.184390 554945.597570 -338546.753770 7
130 CUST6937 52171.166596 220399.339887 -36247.379153 5
153 CUST8155 49397.810737 266000.319070 -139503.197730 5
154 CUST8250 51282.494454 243326.951520 -158946.685580 5
157 CUST8288 59954.562314 346459.485800 -275515.697980 5
162 CUST8461 48295.398164 336834.753390 -336834.753390 5

```

```

BalanceSegment VolumeSegment
1 LowBal HighVol
53 LowBal HighVol
54 LowBal MidVol
80 LowBal HighVol
94 LowBal MidVol
99 LowBal MidVol
101 LowBal HighVol
109 LowBal HighVol
111 LowBal HighVol
130 LowBal MidVol
153 LowBal HighVol
154 LowBal MidVol
157 LowBal HighVol
162 LowBal HighVol

```

#### Task 4: Financial Risk Identification

- Track accounts with frequent large withdrawals or overdrafts.

- Calculate balance volatility using standard deviation or coefficient of variation.

- Use IQR or z-score methods to detect anomalies.

- Highlight customers with irregular or suspicious transaction behavior.

```

In [23]: # step 4.1 • Track accounts with frequent large withdrawals or overdrafts.
withdrawals = df[df["TransactionType"].str.contains("withdraw", na=False)]
threshold = withdrawals["TransactionAmount"].quantile(0.90)

large_withdrawals = withdrawals[withdrawals["TransactionAmount"] >= threshold]

large_withdrawals_by_acct = large_withdrawals.groupby("AccountID").size().reset_index(name="LargeWithdrawalCount")

# Overdraft proxy: negative balance
overdraft_accts = df[df["AccountBalance"] < 0].groupby("AccountID").size().reset_index(name="OverdraftTxnCount")

print(overdraft_accts)
print(large_withdrawals_by_acct)

AccountID OverdraftTxnCount
0 ACC16241 1
1 ACC19178 1
2 ACC23736 1
3 ACC26973 1
4 ACC28154 1
5 ACC28292 2
6 ACC29477 1
7 ACC33287 1
8 ACC49774 1
9 ACC58667 1
10 ACC70314 1
11 ACC77533 1
12 ACC83005 1
13 ACC88449 1
14 ACC94242 1
Empty DataFrame
Columns: [AccountID, LargeWithdrawalCount]
Index: []

```

```

In [24]: # step 4.2 • Calculate balance volatility using standard deviation or coefficient of variation.

bal_vol = df.groupby("AccountID")["AccountBalance"].agg(["mean", "std"]).reset_index()
bal_vol["CV"] = bal_vol["std"] / bal_vol["mean"].replace(0, np.nan)

print(bal_vol)

   AccountID      mean       std        CV
0  ACC10117  70107.007957  25886.972758  0.369249
1  ACC10996  43568.008084  9434.002316  0.216535
2  ACC11062  38137.132610  3208.737888  0.084137
3  ACC11188  69652.151044  35494.660810  0.509599
4  ACC11285  97401.348560  55922.732441  0.574147
.. ...
189 ACC97225  38652.306677  28069.592780  0.726207
190 ACC97411  55978.315635  7871.678922  0.140620
191 ACC99117  47228.185087  20780.582578  0.440004
192 ACC99409  83743.915565  21429.756821  0.255896
193 ACC99549  68641.201433  26251.797058  0.382450
[194 rows x 4 columns]

```

```

In [25]: # step 4.3 • Use IQR or z-score methods to detect anomalies.

# z-score :
import scipy.stats as stats

df['tZScore'] = np.abs(stats.zscore(df['TransactionAmount']))
z_anomalies = df[df['tZScore'] > 3]
print(f"Z-score anomalies (>3): {len(z_anomalies)}")
print(df['tZScore'])

print("---*50)

df['aZScore'] = np.abs(stats.zscore(df['AccountBalance']))
z_anomalies = df[df['aZScore'] > 3]
print(f"Z-score anomalies (>3): {len(z_anomalies)}")
print(df['aZScore'])

```

```
Z-score anomalies (>3): 0
0    1.266887
1    1.148219
2    1.501730
3    0.256222
4    0.355519
...
795   2.431868
796   1.642768
797   0.389231
798   0.969360
799   1.018040
Name: tZScore, Length: 800, dtype: float64
```

```
Z-score anomalies (>3): 3
0    0.048835
1    1.457611
2    0.870504
3    1.228560
4    1.600137
...
795   0.008947
796   0.762771
797   1.094223
798   0.926659
799   0.368842
Name: aZScore, Length: 800, dtype: float64
```

In [26]: # step 4.4 • Highlight customers with irregular or suspicious transaction behavior.

```
customer_stats = (
    df.groupby("CustomerID")["TransactionAmount"]
        .agg(["mean", "std"])
        .reset_index()
)
```

```
print(customer_stats)

CustomerID      mean      std
0   CUST1042  54043.209545  30565.484684
1   CUST1114  48225.320740  26398.893405
2   CUST1121  99551.279378  6457.060342
3   CUST1189  39113.139630  14202.571329
4   CUST1223  33888.232900  21109.744323
...
183  ...  ...
184  ...  ...
185  ...  ...
186  ...  ...
187  ...  ...

[188 rows x 3 columns]
```

In [27]: df = df.drop(columns=["mean", "std"], errors="ignore")
df = df.merge(customer\_stats, on="CustomerID", how="left")
df.tail()

Out[27]:

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product	Firm	Region	Manager	TransactionDate	...	Credit	Debit	Sign	NetAmount	Year	Month	tZScore	aZScore
795	11	CUST8461	ACC60432	Current	Withdrawal	Credit Card	Firm E	Central	Manager 3	2023-02-17	...	0.0	0.0	-1	-119878.30500	2023	2023-02	2.431868	0.008
796	44	CUST1121	ACC41829	Credit	Withdrawal	Personal Loan	Firm C	East	Manager 4	2023-10-26	...	0.0	0.0	-1	-97933.33752	2023	2023-10	1.642768	0.76
797	160	CUST3059	ACC28292	Current	Payment	Mutual Fund	Firm E	South	Manager 4	2023-05-11	...	0.0	0.0	-1	-63072.38174	2023	2023-05	0.389231	1.094
798	37	CUST1042	ACC28295	Current	Payment	Personal Loan	Firm E	West	Manager 3	2023-03-26	...	0.0	0.0	-1	-25289.82472	2023	2023-03	0.969360	0.926
799	101	CUST2464	ACC53865	Loan	Deposit	Mutual Fund	Firm C	South	Manager 4	2023-07-19	...	0.0	0.0	1	23936.04314	2023	2023-07	1.018040	0.368

5 rows × 25 columns

In [28]: df["std"] = df["std"].fillna(0)

In [29]: df["Suspicious\_Flag"] = np.where(
 df["TransactionAmount"] > (df["mean"] + 3 \* df["std"]),
 "Suspicious",
 "Normal"
)

In [30]: suspicious\_customers = df[df["Suspicious\_Flag"] == "Suspicious"][
 ["CustomerID", "TransactionAmount", "mean", "std"]
]
print(suspicious\_customers)

Empty DataFrame  
Columns: [CustomerID, TransactionAmount, mean, std]  
Index: []

Task 5: Visualisation

● Conduct extensive exploratory data analysis with attractive visualizations for your findings

In [31]: # step 5.1 • Conduct extensive exploratory data analysis with attractive visualizations for your findings.

```
## histogram plot of Account balance

sns.set_theme(style="whitegrid")

plt.figure(figsize=(16,9))

ax = sns.histplot(
    df["AccountBalance"],
    bins=30,
    kde=True,
    color="blue"
)

plt.title("Balance Distribution")
plt.xlabel("Account Balance")
plt.ylabel("Count")

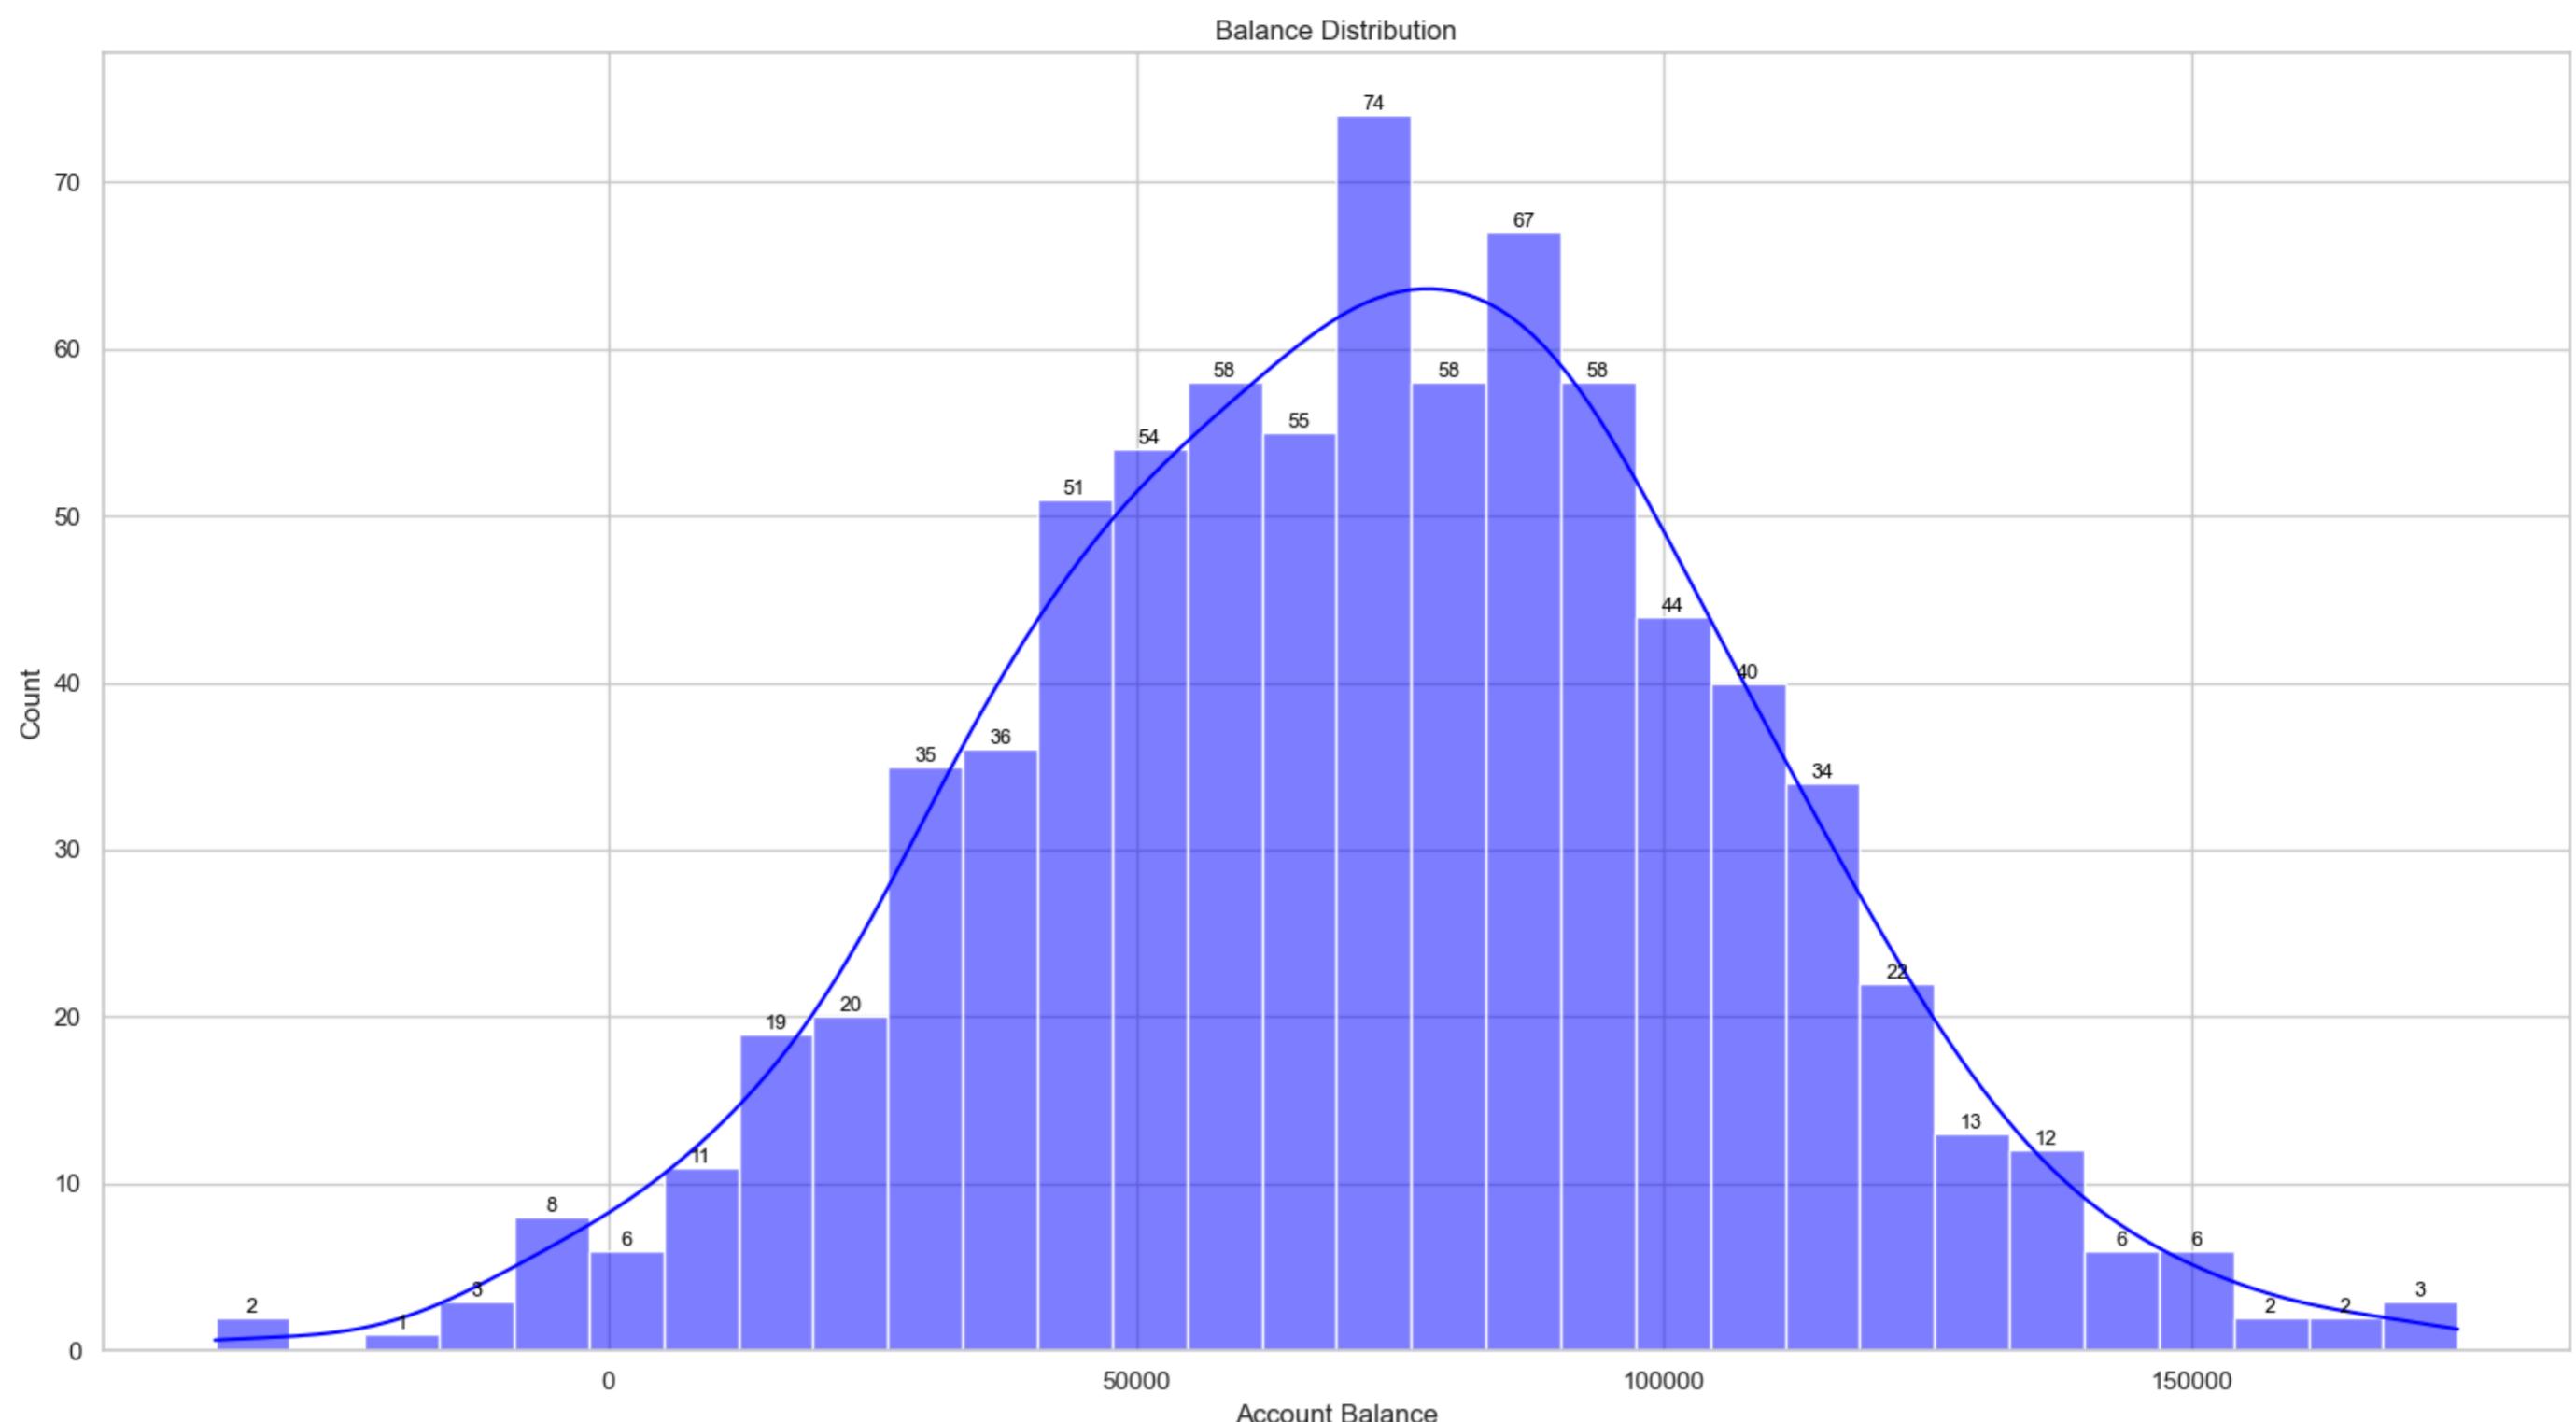
# ---- ADD DATA LABELS ----
for patch in ax.patches:
    height = patch.get_height()
    if height > 0:
        ax.annotate(
            f"{int(height)}",
            (patch.get_x() + patch.get_width() / 2, height),
            ha="center",
            va="bottom",
```

```

        fontsize=9,
        color="black"
    )

plt.tight_layout()
plt.show()

```



In [32]: # a:) Scatter plot of average balance vs total volume :

```

plt.figure(figsize=(16, 9))
sns.scatterplot(data=custom_summary, x='TotalVolume', y='AvgBalance', hue='AvgBalance', palette='viridis')
plt.title('Average Balance vs. Total Transaction Volume')
s=140,          # medium dot size
alpha=0.9
plt.show()

## b:) transaction amount outliers :
import numpy as np

plt.figure(figsize=(12,6))

ax = sns.boxplot(
    y=df["TransactionAmount"],
    color="royalblue"
)

plt.title("Transaction Amount Outliers")
plt.ylabel("Transaction Amount")
plt.xlabel("")

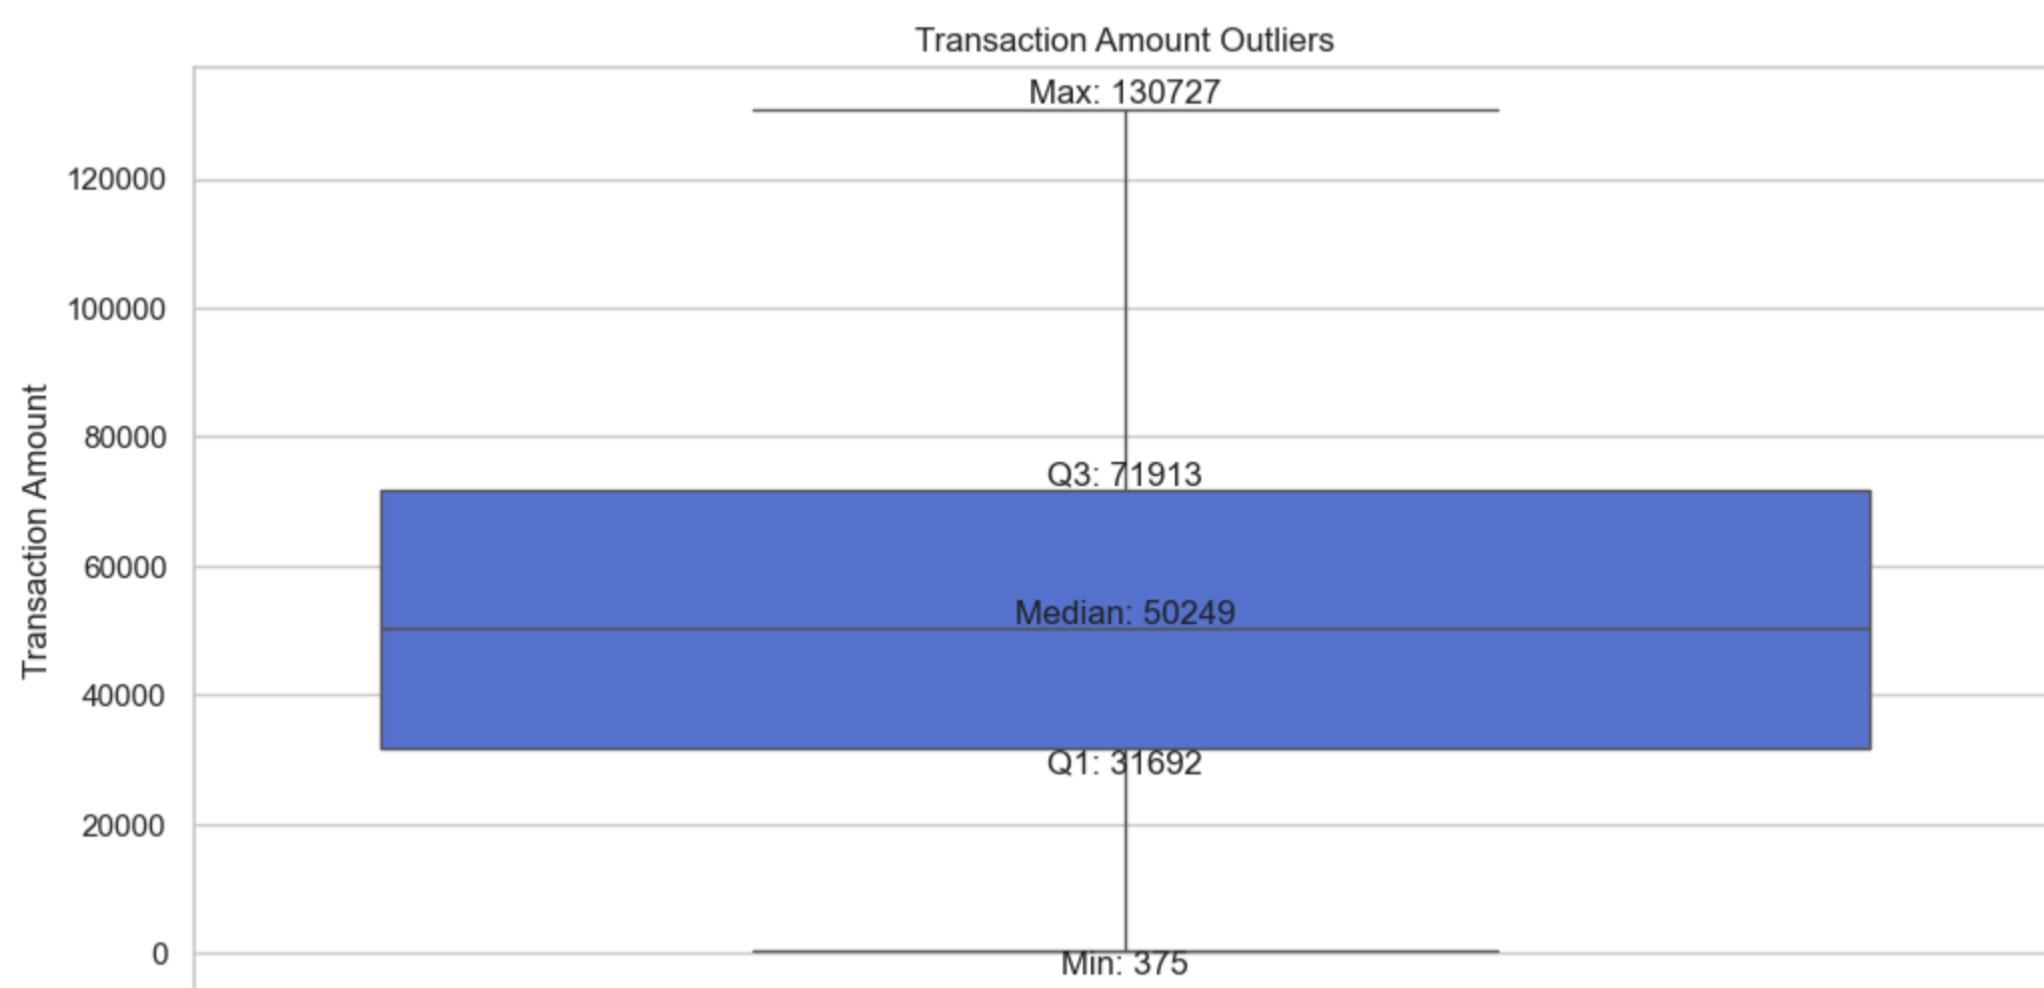
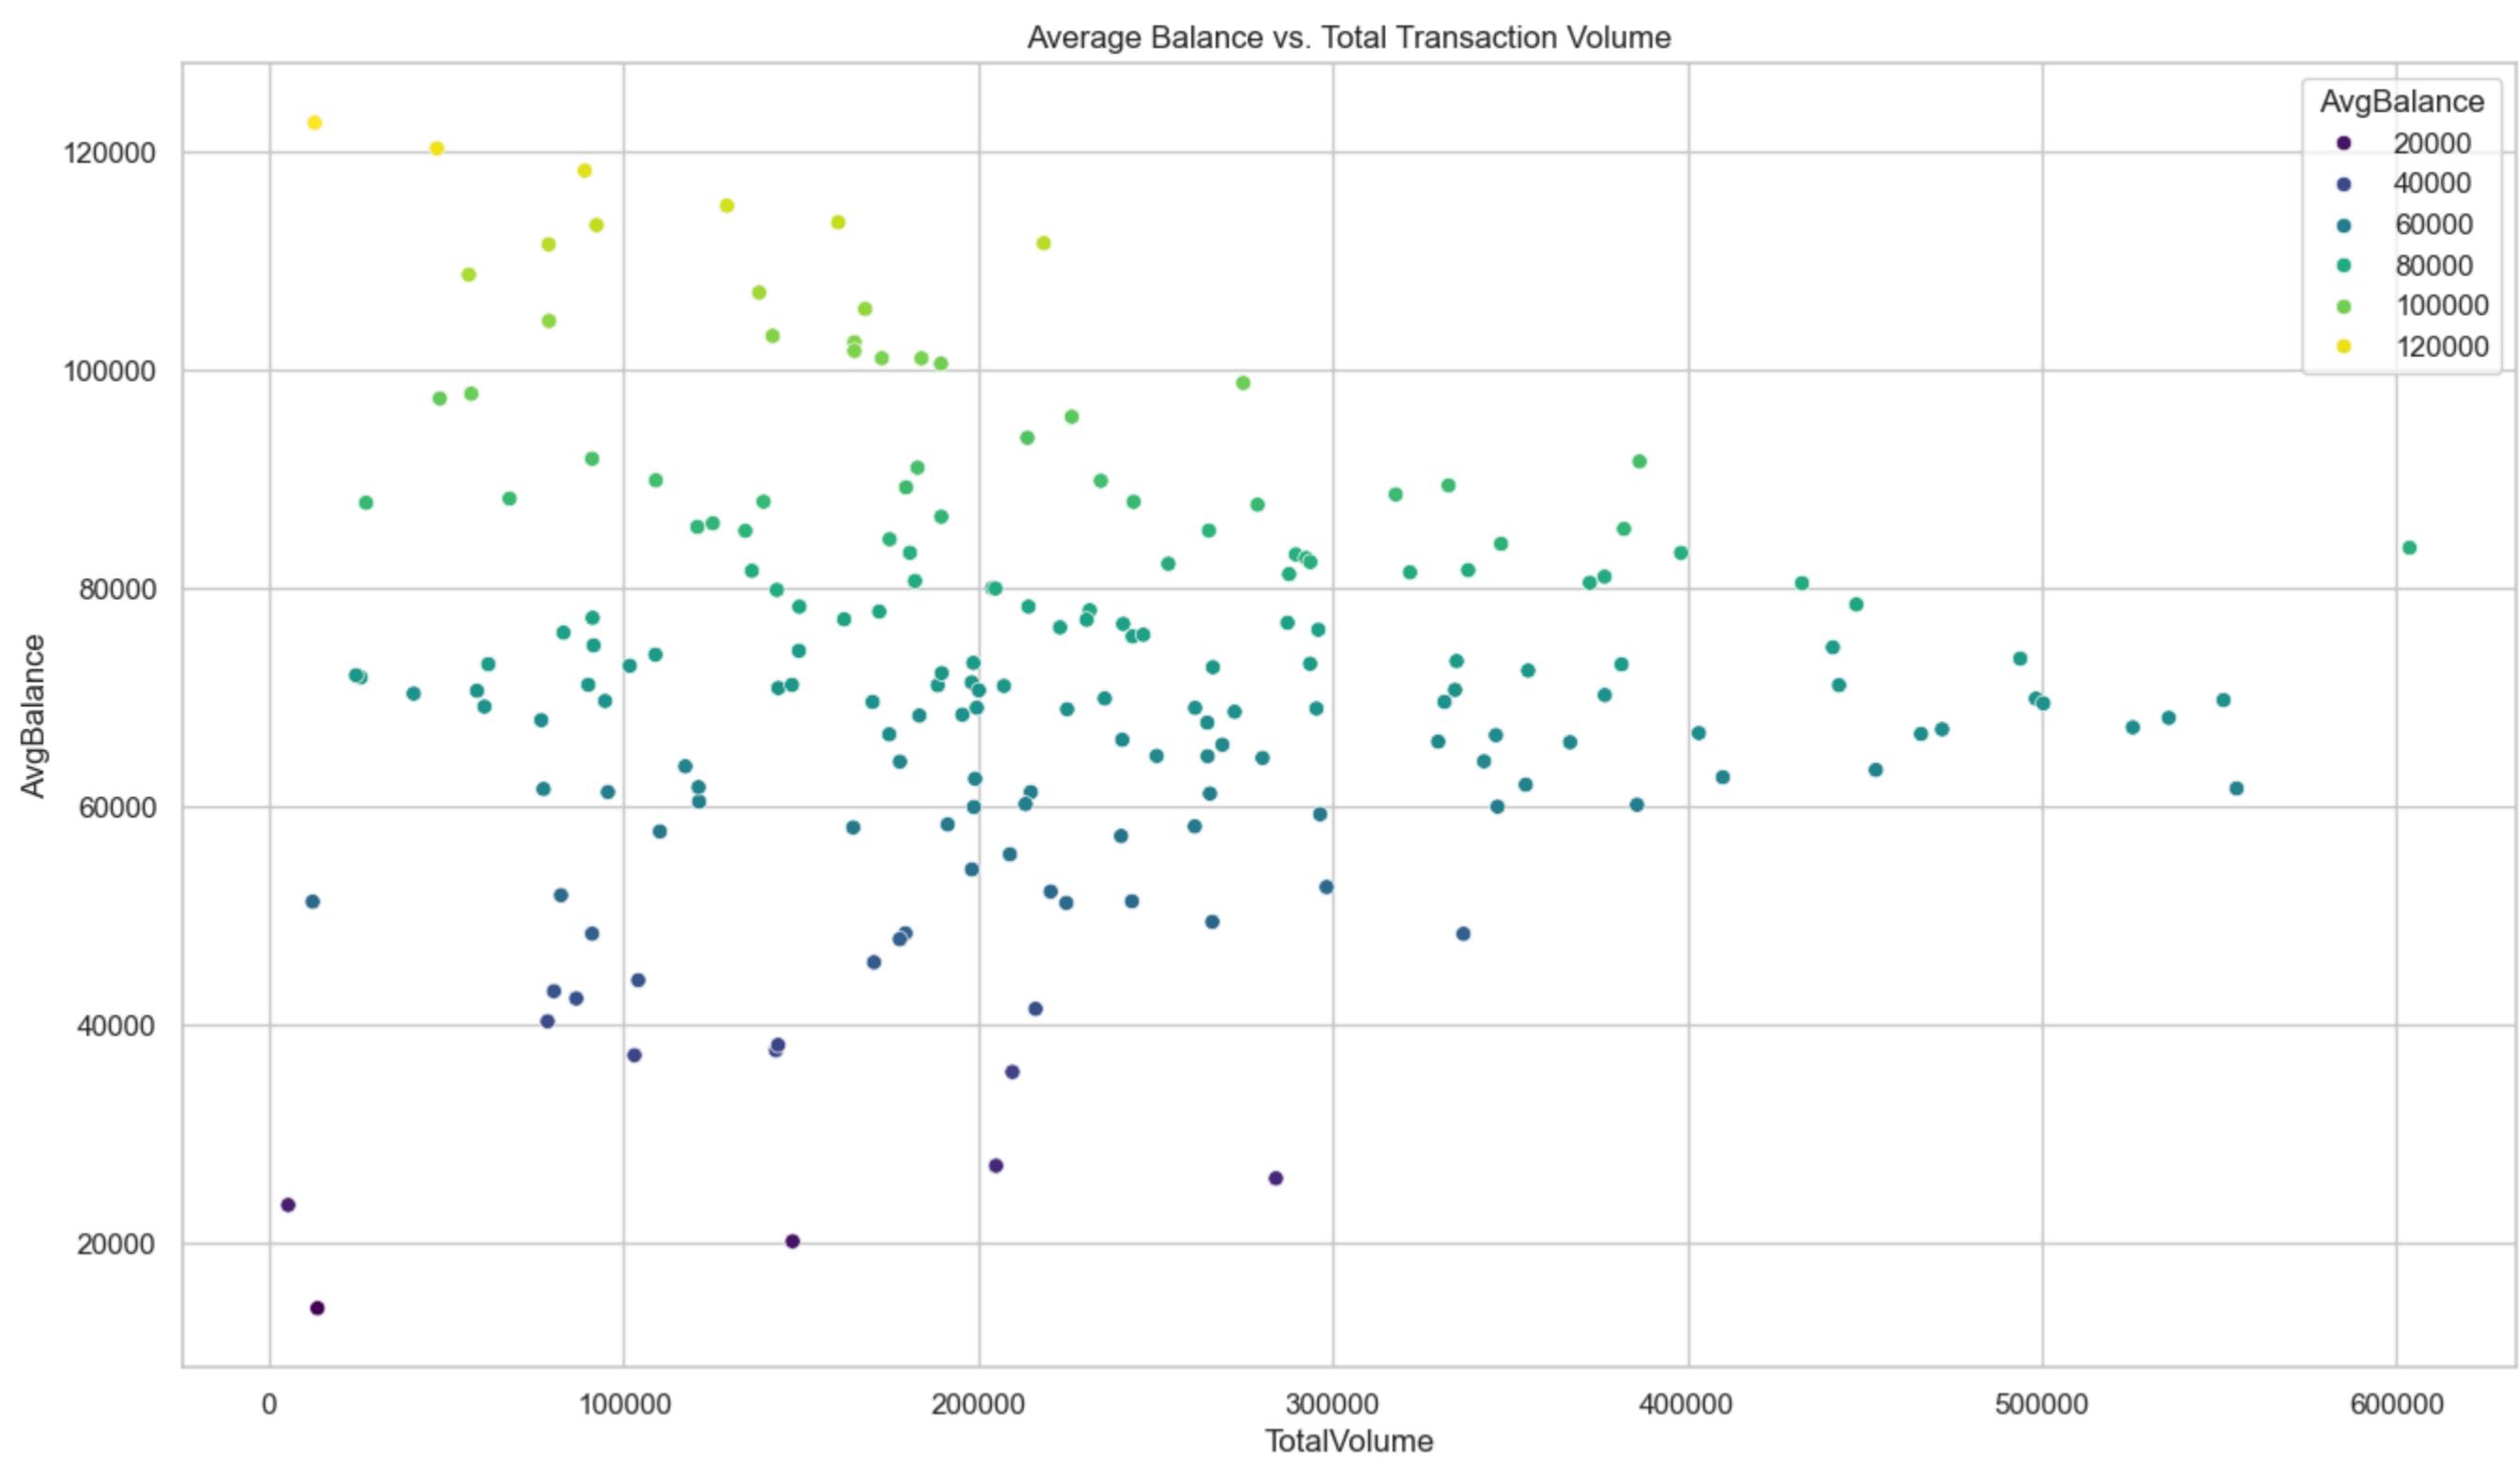
# ---- Calculate statistics ----
data = df["TransactionAmount"].dropna()

q1 = np.percentile(data, 25)
median = np.percentile(data, 50)
q3 = np.percentile(data, 75)
min_val = data.min()
max_val = data.max()

# ---- Add labels ----
ax.text(0, median, f"Median: {median:.0f}", ha="center", va="bottom")
ax.text(0, q1, f"Q1: {q1:.0f}", ha="center", va="top")
ax.text(0, q3, f"Q3: {q3:.0f}", ha="center", va="bottom")
ax.text(0, min_val, f"Min: {min_val:.0f}", ha="center", va="top")
ax.text(0, max_val, f"Max: {max_val:.0f}", ha="center", va="bottom")

plt.show()

```



#### Task 6: Hypothesis Testing

- Test whether high-volume transaction accounts have statistically higher average balances than low-volume accounts.
- Conduct hypothesis testing based on segmentation.

```
In [33]: # step 6.1 ● Test whether high-volume transaction accounts have statistically higher average balances than low-volume accounts.
from scipy import stats

# Segment by volume
q70 = cust_summary["TotalVolume"].quantile(0.70)
q30 = cust_summary["TotalVolume"].quantile(0.30)

high_vol = cust_summary[cust_summary["TotalVolume"] >= q70]["AvgBalance"].dropna()
low_vol = cust_summary[cust_summary["TotalVolume"] <= q30]["AvgBalance"].dropna()

# Welch's t-test
t_stat, p_val = stats.ttest_ind(high_vol, low_vol, equal_var=False)

t_stat, p_val
```

```
Out[33]: (np.float64(-0.386478657678767), np.float64(0.7001935724836169))
```

```
In [34]: # step 6.2 ● Conduct hypothesis testing based on segmentation.
df["AccountBalance"] = pd.to_numeric(df["AccountBalance"], errors="coerce")
txn_count = (
    df.groupby("CustomerID")
    .size()
    .reset_index(name="Transaction_Count")
)
```

```
In [35]: df = df.drop(columns=["Transaction_Count"], errors="ignore")
df = df.merge(txn_count, on="CustomerID", how="left")
```

```
In [36]: median_txn = df["Transaction_Count"].median()

df["Customer_Segment"] = df["Transaction_Count"].apply(
    lambda x: "High Volume" if x > median_txn else "Low Volume"
)
```

```
In [37]: high_volume_bal = df[df["Customer_Segment"] == "High Volume"]["AccountBalance"].dropna()
low_volume_bal = df[df["Customer_Segment"] == "Low Volume"]["AccountBalance"].dropna()
```

```
In [38]: from scipy.stats import ttest_ind

t_stat, p_value = ttest_ind
```

```
    high_volume_bal,  
    low_volume_bal,  
    equal_var=False  
)  
  
t_stat, p_value  
  
if p_value < 0.05:  
    print("Reject Null Hypothesis: Significant difference exists")  
else:  
    print("Fail to Reject Null Hypothesis: No significant difference")
```

Fail to Reject Null Hypothesis: No significant difference

task 7: Video Presentation

- Record a short video summarizing findings and insights.
- Highlight what drives customer transaction behavior and financial risk.
- Discuss data-backed recommendations for customer engagement or monitoring.

First video: [Goldman Sachs Presentation video no. I](#)

Second video: [Goldman Sachs Presentation video no. II](#)